

Replication / Computer Science

[Re] Velho and Legrand (2009) - Accuracy Study and Improvement of Network Simulation in the SimGrid Framework

Arnaud Legrand^{1,2,3,4,5, ID} and Pedro Velho⁶¹Univ. Grenoble Alpes, Grenoble, France – ²CNRS, Délégation Alpes, Grenoble, France – ³INRIA Rhone Alpes, Grenoble, France –⁴Grenoble INP (Institute of Engineering Univ. Grenoble Alpes), Grenoble, France – ⁵Laboratoire d'Informatique de Grenoble, CNRS UMR5217, 38000 Grenoble, France – ⁶Ryax Technologies, Lyon, FranceEdited by
(Editor)Received
01 November 2018Published
–DOI
–

This paper reports the successful reproduction of the results in the article [1] entitled *Accuracy Study and Improvement of Network Simulation in the SimGrid Framework*, which has been published at the SimuTools 2009 conference. We detail several pitfalls we stumbled upon during this reproduction process and report the actions we took to improve the reproducibility of this work.

The first action we took is related to the visibility and availability of the original article. Open access was not mandated by the funders of this research at the time of publication. Even worse, we were regularly told about constraining copyright issues from the editors and discouraged to make our articles publicly available. Only the bibliography entry was available on HAL and not the PDF. Yet the preprint was hosted on the webpage of both authors and visible mostly through the SimGrid publication webpage.

Action #1: The PDF version of the original article has thus been uploaded on HAL when engaging in the replication process.

1 Historical Context

This article compares the accuracy of two methods for predicting how competing TCP network flows interfere with each others. It is the first article Arnaud Legrand wrote with his first PhD student Pedro Velho. It was already a reproduction of the work [2] of close colleagues, Henri Casanova and Kayo Fujiwara, and we already had faced difficulties in doing so at that time. Actually, we could never obtain the exact same numbers as them despite their care and ours. This failure motivated us to improve our methodology, and in particular switching to R and later to org-mode, but it was 10 years ago. Trying to reproduce this article was thus a good test of time!

1.1 Scientific context

SimGrid is a simulation toolkit designed to help with the performance evaluation of large scale distributed computing systems such as data/desktop grids, clusters or peer-to-peer systems. In this field, it is common to resort to simulation which enables to explore many application and platform scenarios, including platforms which do not exist yet, and to obtain reproducible results (in theory). Unfortunately, as noted in [3], most researchers build their own *ad hoc* simulators which are rarely made (1) available

Copyright © 2020 A. Legrand and P. Velho, released under a Creative Commons Attribution 4.0 International license.

Correspondence should be addressed to Arnaud Legrand (arnaud.legrand@imag.fr)

The authors have declared that no competing interests exist.

Code is available at <https://github.com/alegrand/reproducibility-challenge..>

, which hinders both reproducibility of results and comparison of articles with one another, nor (2) validated, which makes final results quite questionable. SimGrid is a 20-years old attempt to provide a high quality simulation toolkit, which would be (1) stable and perennial from a software point of view and (2) whose models would be as validated as possible against reality and other simulators.

In this context, network simulation is certainly the most critical part and packet-level simulation is often considered as particularly realistic and faithful since they it allows to account for every detail of network protocols. Unfortunately, such *microscopic* approach is undoubtedly interesting when studying peculiarities of network protocols but it leads to prohibitively long simulation time when studying large-scale distributed systems. An alternative is thus to simulate networks by relying on higher level *macroscopic* models, thus enabling much faster simulation at the potential cost of an accuracy loss. SimGrid, uses a flow-level approach that approximates the behavior of TCP networks, including TCP's bandwidth sharing properties. A preliminary study of the accuracy loss by comparing it to popular packet-level simulators has been proposed by Fujiwara and Casanova in [2] and in which regimes in which SimGrid's accuracy was comparable to that of these packet-level simulators were identified. The article we reproduce here [1] was a reproduction these experiments and provided a deeper analysis that enabled to greatly improve SimGrid's range of validity.

The interconnection network is modeled as a graph where nodes represent hosts while edges represent network links. In SimGrid's flow-level modeling, the time needed to transfer a message of size S between hosts i and j is given by:

$$T_{i,j}(S) = L_{i,j} + S/B_{i,j}, \quad (1)$$

where $L_{i,j}$ (resp. $B_{i,j}$) is the end-to-end network latency (resp. bandwidth) on the route connecting i and j . Although determining $L_{i,j}$ may be straightforward, estimating the bandwidth $B_{i,j}$ is more difficult as it depends on interactions with every other flow. This is generally done by assuming that the flow has reached *steady-state*, in which case the simulation amounts to solving a bandwidth sharing problem, i.e., determining how much bandwidth is allocated to each flow.

More formally, consider a connected network that consists of a set of links \mathcal{L} , in which each link l has capacity B_l . Consider a set of flows \mathcal{F} , where each flow is a communication between two network vertices along a given path. Determine a "realistic" bandwidth allocation ρ_f for flow f , so that:

$$\forall l \in \mathcal{L}, \quad \sum_{f \text{ going through } l} \rho_f \leq B_l. \quad (2)$$

In SimGrid, the "realistic" bandwidth sharing model [4] used is Max-Min fairness [5], which is reached by recursively maximizing

$$\min_{f \in \mathcal{F}} w_f \rho_f \quad \text{under constraints in Eq. (2),} \quad (3)$$

where w_f is generally chosen as the round-trip time of flow f . This objective corresponds to what one would naively expect from a network, i.e. be "as fair as possible" so that the least favored flows receive as much bandwidth as possible while accounting through weights w_f for the well-known RTT-unfairness of TCP [6].

Given the computed bandwidth allocation (which defines all data transfer rates), and the size of the data to be transmitted by each flow, one can determine which flow will complete first. Upon completion of a flow, or upon arrival of a new flow, the bandwidth allocation can be reevaluated. Compared to a packet-level simulation, this approach allows to quickly step-forward in time when large data transfers are involved. However, since steady-state is assumed, it ignores many transient aspects such as throughput oscillations and slow start. This (in)validation work was later extended to compare with

other bandwidth sharing models [7] and has been the core of the PhD thesis of Pedro Velho [8].

In the article we reproduce, the accuracy of the *flow-level* simulations of SimGrid are compared to the *packet-level* simulations of GTNetS [9] the Georgia Tech Network Simulator. This was done through three series of simulations

1. One-link: The first set of experiments is for a single TCP flow going through a single link with varying physical **latency** and **bandwidth**, and message **size**. The main goal of this scenario was to study the size for which transient effects such as slow start are negligible.
2. A Dumbbell Topology: The second set of experiments is for two TCP flows A and B on a dumbbell topology with varying **bandwidth** of the inner link and **latency** of the end-link used by flow B. The main goal of this scenario was to study the ability of accounting for RTT-unfairness.
3. Random Topology: 4 sets of 10 random topologies generated with a topology generator were used. The sets comprised either *small* (50 nodes) or *large* (200 nodes) and either relatively *homogeneous* or *heterogeneous* platforms. 200 flows were generated between random pairs of end-points in the topology, which all start simultaneously and communicate 100MB of data. The main goal of this scenario was to evaluate the overall accuracy of SimGrid and possibly to detect corner-case situations for which the SimGrid model was particularly wrong.

Due to the long simulation time, we only reproduce in this article the first series of simulation but we checked that we could easily run at least one simulation of the two other series.

1.2 Computational context

SimGrid is mostly written in C while GTNetS is mostly written in C++ and both are open source simulators. Although SimGrid is designed to be as stand alone as possible, GTNetS relies on third party libraries. **The first challenge would thus be to reproduce a software environment allowing to recompile and rerun both libraries.**

To ease the comparison of both simulators, SimGrid had been modified to run GTNetS internally, which allowed to easily switch between the microscopic (GTNetS) model and the macroscopic (Max-Min) model from the command line, while using the exact same platform description and communication scenario. This integration required modifying both SimGrid and GTNetS and was done through a set of patches before being partly integrated in the main branch of SimGrid. **The second challenge would thus be to manage to correctly modify and recompile a simulator using both libraries.**




Figure 1. The simulation workflow

Although these details were not given in the articles, it could be recovered from one of the README we found (see Figure 3 and Figure 4) and the general workflow of the simulations for all three scenarios was as follow (see Figure 1):

- A simple C code called `gtnets.c` was linked against SimGrid and GTNetS;

- A perl script called `sweep-parse.pl` (when called with the `sweep` argument) would generate platform and flow/deployment XML input files and run all simulations by passing the previous XML input file to the `gtnets` binary with a different command line argument to switch between the GTNetS model and the Max-Min model. The simulation would produce a text output.
- The same perl script (when called with the `parse` argument) would then parse all the text logs and produce a csv data file.
- The data file would then be analyzed with an R script and since our mastery of R was quite low at that time, we still relied on gnuplot to generate figures.

The third challenge would thus be to manage to run all this workflow, provided the right instructions could be found.

Note that although the first two series of experiments did not have much external dependencies, the third one relied on many random network topologies generated by BRITE [10], which is a discontinued Java software, using the Waxman model [11]. The description of the parameters used to generate the topologies were shallow and there was no information regarding seeds so our hope to rerun this software to regenerate the same topologies was quite low. However, these intermediate files may have been stored and made available. **The fourth challenge would thus be to recover the network topology and data used in the third series of experiments.**

2 Rebuilding the code and its environment

2.1 Original source code and retrieval of the software

Instructions — Although the development of SimGrid is still very active, GTNetS' development appears to be discontinued as the last version of GTNetS dates back October 2008. Finding both source code is relatively easy however, the main difficulty was to find the instructions and to know which version to use. SimGrid has successively moved from the Inria gforge to the Inria gitlab, GitHub, and more recently Framagit. Although the whole software history has been correctly moved in the process, we realized some information have not been transferred and even sometimes lost:

- Although we could have used a development version of SimGrid from late 2009, we thought it would be simpler to reproduce this work using a stable release which integrates the GTNetS support (e.g., the version 3.3, which dates from April 2009). Unfortunately, the releases of SimGrid on GitHub only start from May 2010. Indeed, although the SimGrid project started in 2000, its git history only starts in 2004 as the CVS history was not migrated to SVN (it was considered of little interest). Later, when the development team decided to migrate from SVN to git (in 2010), the SVN tags have not been transferred. Likewise, when the project migrated from the Inria gforge, not all releases (as an archive) of SimGrid were transferred. Fortunately, all the old releases of SimGrid are still available on the Inria gforge.

Action #2: We have thus now uploaded the original release of SimGrid version 3.3 on Github.

This upload is mostly manual and will be done for other old releases as soon as possible.

- The \LaTeX source of the article is stored in the private Inria Gforge `simgrid-publics` project, in an `svn` under the `PUBLISHED/09_validation_simutools` directory.

Action #3: We have now made the \LaTeX source of the article available in the github repository attached to this Rescience submission.

- We remembered that we had made our instructions on how conduct these experiments available somewhere but no link was given in the original article and we could not really remember where it was as there was no standard way of doing so back then. We initially thought they were given on the former `contrib/` section of <http://simgrid.gforge.inria.fr/> (, which was hard to maintain and was thus abandoned) or on <http://simgrid-publis.gforge.inria.fr/> (, which finally only hosts data on two articles from 2011). However even after inspecting the Internet Archive, we could not find it.

Arnaud Legrand therefore tried to find them on his laptop but although he could find many related files (including the topology generators) he failed finding the right data and doing so, he realized many the data of some of his previously published articles were dangling links and had not been correctly transferred when migrating from a laptop to an other! The instructions could probably have been recovered on his backup hard drives but he had the chance to meet Pedro Velho and to ask him whether his own backups were in better shape, which was fortunately the case. Pedro Velho could find all the required data (a 61MB zip archive) and shared it with him. It turned out that we later realized that this archive was also simply available from Pedro Velho's former webpage, which is still available but which is not highly ranked on search engines and which he cannot modify anymore as he now works for a different company.

Action #4: We have now made the instructions and data used in the original article available in the GitHub repository attached to this Rescience submission.

This archive comprises 3 sub-archives corresponding to each of the 3 series of simulations mentioned earlier (01-oneLink.tgz, 02-dumbbell.tgz, 03-random.tgz) as well as a GTNetS version (GTNetS-Oct-10-08.tar.gz) and the master simulation file (gtnets.c) which should be compiled against SimGrid and GTNetS. The README that can be found in each sub-archive (see Figure 4) describes in details how to rerun the experiments and corresponds to the process described in Section 1.2. A good surprise was that the third archive contained all the random graphs used in the simulation, hence saving us the burden of trying to regenerate them with BRITE. Unfortunately, the master README (see Figure 3) provides information about dates and the contents of the archive but most information related to software versions are broken (it was a working version, which we intended to consolidate when the article would be accepted). Furthermore, after having compiled GTNetS and a thorough inspection of the source code, we realized it did not seem to have been modified to work with SimGrid.

- Arnaud Legrand therefore started searching again for GTNetS versions that would be on his laptop and finally found one, along with all the patches and compiling instructions which are crucial to correctly build such prototype software (see Figure 5). These information were actually public but had become completely hidden in the (now unmaintained and long forgotten) `contrib` section of the SVN (while git is now the default version manager) of the Inria Gforge SimGrid project.

Action #5: We have now ensured that the GTNetS version and the patches we used are archived on Software Heritage.^a

^aThe save request was done on 4/30/2020, 6:50:02 PM but it is still pending.

In the end, we have thus managed to recover three important archives (see Figure 2), whose versions should be the one run to produce the results of the original SimuTools 2009 article:

1. The stable release v3.3 of SimGrid (from April 2009) from the public Inria Gforge. Although experiments were probably run in late 2008, the previous stable SimGrid release is from 2007 and v3.3 incorporates everything that was needed.




Figure 2. The three archives required to reproduce this work.

2. A snapshot of GTNetS from January 2008 along with the patches to apply from the public Inria Gforge SimGrid project but which was not visible anymore.
3. The simulation instructions and data, from one of the author's hard drive although they were also available from his website but both authors had forgotten about it.

Almost no information regarding the software environment was available except that it was run on a Debian in the late 2008 (see Figure 5).

2.2 Rebuilding the software environment

SimGrid is mostly a C library whose software dependencies had at that time been kept to the bare minimum (C and C++ compiler). Furthermore, we are developers of the SimGrid library so building it was rather straightforward even on a recent Linux distribution. However, after trying to compile GTNetS, we quickly realized it depends on the Qt3 GUI Library whereas the version which is now commonly found is Qt5! Therefore, we decided to recreate a minimal software environment as close as possible to the one of 2008.

The code name for the stable Debian distribution at that time was *Lenny* Debian provides two particularly interesting tools to reproduce "old" environments:

1. The Debian snapshot archive is a wayback machine that allows access to old packages based on dates and version numbers. It consists of all past and current packages the Debian distribution ever provided.
2. The Debuerreotype is a reproducible, snapshot-based Debian rootfs builder. It allows to prepare old Debian images from the snapshot archive, which is particularly useful to prepare Docker images containing old software environments.

Both authors regularly used `testing` so after investigating a bit on the snapshot archive which versions of the libraries were available and when they have been introduced, we decided to try to bootstrap a debian Lenny from the 1st of May 2009 with the following command:

```
1 debuerreotype-init --keyring=/usr/share/keyrings/debian-archive-removed-keys.gpg \
2 rootfs testing 2009-05-01-T03:27:08Z
```

Building such an image involves installing (with `dpkg`) old packages in a sub-directory pretending you are root. The `keyring` argument passed to `debuerreotype-init` allows to indicate `dpkg` that it is safe to install these old packages even if they have been signed by package maintainers which are currently not active anymore. Unfortunately, although this approach worked like a charm for more recent target dates (e.g., 2015-06-04-T10:47:50Z), it miserably fails with a "Segmentation fault" when installing `base-passwd`:

```
W: Failure trying to run: chroot "/home/alegrand/Work/Documents/Articles/2020/
reproducibility_challenge/simgrid3.3_gtnets/rootfs" dpkg --force-depends
--install /var/cache/apt/archives/base-passwd_3.5.21_amd64.deb
W: See /home/alegrand/Work/Documents/Articles/2020/reproducibility_challenge/
simgrid3.3_gtnets/rootfs/debootstrap/debootstrap.log for details
```

simutools09/instructions/README

Author : Pedro Velho
last modified : 03/11/2008

Disclaimer
#####

This text as well as the data and results provided here are under GPL copyright. To consult the GPL terms and usage condition see in the top directory: GPL.txt

All programs use the gtnets.c simulation program. A source code copy is located in this directory.

CAUTION: This script relies on parsing the output, so every modification (even slight ones) on gtnets.c output may cause the parsing feature to unpredicted behavior. If you are not sure about your gtnets.c file please use the one distributed here.

Short History
#####

This directory hold experiment comparing the SimGrid framework network simulation engine with GTNets. GTNets is a packet level network simulator and we believe it can provide realistic transmission time prediction due to its characteristic of simulating through discrete events the entire TCP protocol stack. This work is the normal continuation of the work presented by Kayo Fujiwara and Henry Casanova in [1] and was submitted to the SimuTools09 conference which is still to be evaluated and accepted/reject.

Directory Structure
#####

Three category of experiments were performed. Each one was tackled separately and are organize in distinct directories as follows:

- * ./01-onelink - Verify message size communication time correlation
- * ./02-dumbbell - Bandwidth sharing experiments
- * ./03-random - Complex platforms to assure model improvements

Global System Requirements
#####

Experiments and analyze scripts are provided to reproduce the graphs presented in [FIXME(Rapport de recherche ou reference simutools)]. To run experiments some minimum system requirements are needed:

- * GTNets patched simgrid version, we kindly provide GTNets with patches in the simgrid contrib svn repository [FIXME]
- * SimGrid, configured and compiled with GTNets support [FIXME]

For plotting graphs and explore the data:

- * R - the gnu version of S [FIXME]
- * Gnuplot [FIXME]

Refernces
#####

[FIXME] The R (GNU S) language website. FIXME URL
[FIXME] Plotting scientific data with Gnuplot. FIXME URL
[FIXME] SimGrid website. FIXME URL
[FIXME] Kayo Fujiwara and Henry Casanova FIXME

Figure 3. The README which accompanies instructions recovered from Pedro Velho on the simulation workflow are very helpful to understand the general process but lack important version information.

simutools09/instructions/01-oneLink/README

Author : Pedro Velho
last modified : 25/11/2008

Disclaimer
#####

This text as well as the data and results provided here are under GPL copyright. To consult the GPL terms and usage condition see in the top directory: GPL.txt

Short history
#####

Experiments with one link are important to validate the linear model. We intend by linear model the assumption that transmission time is correlated to size, bandwidth and latency in some way such as $T=S/B+L$, for instance. More detailed description about this assumption and the evolution of the SimGrid network engine is presented in[FIXME].

Directory structure
#####

This directory contain many files so they are organized in a directory structure as follows:

- * ./log - Tons of output generated by simgrid while running experiments
- * ./dat - The raw.dat file in an R input format
- * ./tmp - Temporary files, such as those generated by R when gnuplot is called within R
- * ./bin - Auxiliary scripts are stored here
- * ./fig - EPS images are generated inside this directory

Running Experiments
#####

- ./bin/sweep-parse.pl
#####

File ./bin/sweep-parse.pl is the most important script it can run the entire set of experiments using or not a grid/cluster infrastructure to improve simulation speed. Two parameters are used, <first-task> <last-task>. Hence this script run all experiments from first-task up to last-task including those number passed as parameters. Before running this script is important to correctly set the working directory as your SimGrid gtnets binary location. All programs here use the gtnets.c program this programs source is normally located in the experiments top directory.

After all log files are corrected collected in directory ./log this is script may be used to parse results generating ./dat/raw.data which will contain all experiments in a R table format.

CAUTION: This script relies on parsing the output, so every modification (even slight ones) on the program output may cause the parsing feature to unpredicted behavior. A gtnets.c version is stored in the top level directory of experiments, this is a trusted version, if you are not sure about your gtnets.c file please use the one distributed here.

- Example: Be aware that executing the script should overwrite stored output.

./bin/sweep-parse.pl sweep 1 1

This generate one output trace file for the first bandwidth parameter in ./log/trace-file-1-1.log

./bin/sweep-parse.pl parse

To parse the output, after all log files have been collected.

- analyze.R
#####

This file contain all the R function used to analyze the data in ./dat/raw.data. The ./dat/raw.data file is generated by ./bin/sweep-parse.pl script as described before.

I'm used to analyze my data within emacs using ESS (Emacs Speaks Statistic) package.

Figure 4. The README which ships with the first set of experiments (01-oneLink.tgz).

simutools09/README.patching_GTNetS

```
=====
GTNetS/Simgrid patch README
author: Pedro Velho
=====
```

Note About this Patch

This patch is intended to work only with GTNetS downloaded from the GTNetS website link:
<http://www.ece.gatech.edu/research/labs/MANIACS/GTNetS/software/gtnets-current.zip> The last time this patch was downloaded was June 12 2008, Seems they don't have much control about new GTNetS and some other flaws it is difficult to precise a verison number.

Getting GTNetS

Two ways of getting GTNetS, one from the gtnets website and svn simgrid contrib projec tree (RECOMMENDED).

```
$ wget http://www.ece.gatech.edu/research/labs/MANIACS/GTNetS/software/gtnets-current.zip
or
$ svn checkout svn://scm.gforge.inria.fr/svn/simgrid/contrib/trunk/GTNetS/
```

Applying the PATCH

```
-----
$ unzip gtnets-current.zip
$ tar zxvf gtnets-current-patch.tgz
$ cat *.patch | patch -p1
```

Compilling GTNetS

```
-----
Enter directory
$ cd gtnets-current
GTNetS is not a very active project for the moment and the portability is really limited.
For the moment we tried out this patch only in linux platforms using:
```

```
gcc (GCC) 4.1.3 20070629 (prerelease) (Debian 4.1.2-13) Linux 2.6.21-2-686
```

Create a Makefile.linux symbolic link

```
$ ln -sf Makefile.linux Makefile
```

Create dependencies list

```
$ make depend
```

To compile debug version

```
$ make debug
```

To compile optimized version

```
$ make opt
```

=== WARNING ===

A lot of warnings are expected but the application should compile just fine. If the makefile insists in compiling some QT libraries please try a make clean before asking for help.

Installing GTNetS

Commands make debug and opt generates respectively libgtsim-opt.so or libgtsim-debug.so. You will need to link ONLY ONE of these libraries using the symbolic link name libgtnets.so, for instance to libgtsim-debug.so:

```
# ln -sf libgtsim-debug.so /<userhome>/usr/lib/libgtnets.so
```

'libsimgid.so' is the name simgrid is configured to search when running ./configure script.

Now just put the library somewhere you know ldd is searching (tip: export LD_LIBRARY_PATH=\$LD_LIBRARY_PATH:/<userhome>/usr/lib/libgtnets.so && sudo ldconfig)

The gtnets source headers are necessary. So you need to copy all headers to a place where your compiler can find them, such as:

```
$ mkdir /<userhome>/usr/include/gtnets
```

```
$ cp -fr SRC/*.h /<userhome>/usr/include/gtnets
```

Compilling SimGrid with GTNetS

Just add the following option when running configure --with-gtnets=/<userhome>/usr

Bug reports, comments, suggestions: pedro.velho@imag.fr

AMD64 bit patch

Some users experienced some problems during compilation on AMD64 bit architecture. We compiled succesfully the gtnets-current package in an:

```
model name      : AMD Opteron(tm) Processor 248
stepping        : 8
cpu MHz         : 2193.160
cache size      : 1024 KB
clflush size    : 64
cache_alignment : 64
Using gcc (GCC) 4.2.3 (Debian 4.2.3-3)
```

We provide a simple patch to do this: AMD64-FATAL-Removed-DUL_SIZE_DIFF-Added-fPIC-compillin.patch

Figure 5. The README which accompanies the GTNetS patches provides many critical information on how to compile GTNetS and SimGrid.

```
error: 'debootstrap' failed!
```

We then decided to cry for help and asked two Debian guru friends, Vincent Danjean and Samuel Thibault. Samuel Thibault indicated us that he had investigated this by using the simpler following command:

```
1 debootstrap wheezy myroot http://archive.debian.org/debian/ shell
```

and that the error message was then slightly more visible

```
dpkg: warning: parsing file '/var/lib/dpkg/status' near line 5 package 'dpkg':
missing description

Package: dpkg
Status: install ok installed
Maintainer: unknown
Version: 1.16.18
```

The problem actually comes from dpkg. When bootstrapping such an image, we try to use old debian packages with a recent dpkg (the one running on our machine) so it is not surprising that it may break. Although it is not the case, the internal format of Debian packages could have evolved and may not be supported anymore with recent versions of dpkg. Likewise, it is somehow a matter of luck that an old binary still works with a recent kernel... Indeed, when using docker or similar container-based approach, we only divert syscalls so if the ABI of the Linux kernel had changed in the meantime, binary codes would simply fail to run. Fortunately, such changes are quite rare and the Linux/Debian community is making incredible efforts to provide super stable backward compatible software so what could be the reason behind this failure?

Surprisingly Vincent Danjean reported me that the command worked like a charm for him, which means some local configuration from his machine could change this behavior. We could actually track back the problem to an ABI modification of the kernel. As explained for example on the Einstein@Home forum, *"On latest Linux distros, vsyscall is defaulted to none for security reasons. However, this breaks some very old binaries, including some binaries from this project that are statically-linked against ancient versions of glibc"*. Vincent had activated this flag a long time ago to run some old proprietary code. Booting the machine while adding `vsyscall=emulate` to the kernel command line solved the problem and allows debuerreotype to build the desired rootfs.

Since this is a bit far-fetched, we decided to trade precision for simplicity by checking whether ready-to-use Docker images were available on the Docker Hub, which is the case.

```
1 docker search debian-lenny shell
```

NAME	DESCRIPTION	STARS
pblaszczyk/debian-lenny	5.0.10 amd64	3
lpenz/debian-lenny-amd64	Debian 5.0.10 Released 10 March 2012 for amd...	1
lpenz/debian-lenny-i386	Debian 5.0.10 Released 10 March 2012 for i386	1
lpenz/debian-lenny-amd64-minbase	Debian 5.0.10 Released 10 March 2012 for amd...	0
...		

Note that the first version of Lenny (5.0.0) was introduced in February 2009 whereas the one easily found on the DockerHub is the last version (5.0.10) which dates from March 2012. The main differences are related to security updates and should be of not importance for our concern. We arbitrarily chose the `lpenz/debian-lenny-i386` one but according to the instructions of Figure 5 `lpenz/debian-lenny-amd64` should have worked as well. After playing a bit interactively in this Docker image trying to install everything we needed to build GTNetS and SimGrid, and following the patching and compiling instructions, we ended up with the Dockerfile presented in Figure 6. The image can be simply built with the following command:

```

1 docker build -t alegrand/simgrid3_3_gtnets simgrid3.3_gtnets
_____ shell _____
_____ simutools09/simgrid3.3_gtnets/Dockerfile _____

FROM lpenz/debian-lenny-i386

LABEL maintainer="Arnaud Legrand <arnaud.legrand@imag.fr>"

# Software dependencies
RUN apt-get update \
    && apt-get install -y --force-yes gcc g++ make wget unzip subversion patch \
    less libqt3-mt libqt3-headers libqt3-mt-dev qt3-dev-tools

# Code downloading and assembly dependencies
RUN apt-get update \
    && apt-get install -y --force-yes wget unzip subversion patch less

# Downloading GTNetS
RUN cd /root; svn checkout svn://scm.gforge.inria.fr/svn/simgrid/contrib/trunk/GTNetS/
# Downloading SimGrid
RUN cd /root; wget https://gforge.inria.fr/frs/download.php/file/21430/simgrid-3.3.tar.gz

# Building GTNetS
RUN cd /root/GTNetS; unzip gtnets-current.zip ; tar zxvf gtnets-current-patch.tgz
RUN cd /root/GTNetS/gtnets-current; cat ../00*.patch | patch -p1
RUN cd /root/GTNetS/gtnets-current; ln -sf Makefile.linux Makefile && make depend && make opt

# Installing GTNetS
RUN cd /root/GTNetS/gtnets-current/ && \
    mkdir -p /root/usr/lib/ && \
    ln -sf `pwd`/libgtsim-opt.so /root/usr/lib/libgtnets.so && \
    ln -sf `pwd`/libgtsim-opt.so /usr/lib/libgtnets.so && \
    mkdir -p /root/usr/include/ && \
    cp -fr SRC/*.h /root/usr/include/

# Building SimGrid
RUN cd /root/ && tar xzf simgrid-3.3.tar.gz
RUN cd /root/simgrid-3.3/ && \
    ./configure --with-gtnets=/root/usr/ && \
    export LD_LIBRARY_PATH=/root/usr/lib/libgtnets.so && \
    ldconfig && \
    make

RUN apt-get clean

```

Figure 6. The Dockerfile recipe which allows to build both GTNetS and Simgrid

Action #6: We have now proposed a simple and automated way to build a minimal-ist environment comprising the simulation code used in the original article. This Dockerfile recipe has been made available in the GitHub repository attached to this Rescience submission. The resulting docker image has been made available on the DockerHub. It can be recovered using:

```

1 docker pull alegrand38/simgrid3_3_gtnets
_____ shell _____

```

Note that, as such, this Dockerfile is still a bit fragile as it depends on a third party base image (lpenz/debian-lenny-i386) and downloads the code from gforge.inria.fr. Ideally, it would be improved to build on my own debuerreotype image for a specific date and to download the code from [software heritage](https://software-heritage.org/). We propose to leave this for the next reproducibility challenge in a decade or so.

3 Execution and reproduction of results

3.1 Expectations

Following the information from the README of each series of simulations (see Figure 4), we could easily determine which scripts to run (`sweep-parse.pl`). It is interesting to note that the logs of each simulation were stored in the archive (in `log/`) as well as the parsing of these logs (in `dat/`).

Before trying to rerun all this, we ensured a specific parameter combination could be run to manually check whether outputs are matching or not. Here was the target:

```
1 head -4 simutools09/instructions/01-onelink/dat/raw.data
```

```
Bandwidth Latency Size Model Time
1 1.000000e+05 0.00001 1000 CM02 0.010010
2 1.000000e+05 0.00001 1000 GTNets 0.013140
3 1.000000e+05 0.00001 1000 LegrandVelho 0.010974
```

And here was the output we should get from running `gtnets`.

```
1 head -46 simutools09/instructions/01-onelink/log/trace-file-1-1.log
```

```
>=====<
=====> Bandwidth (B) : 1.000000e+05 B/s (Bytes per second)
=====> Latency (L) : 0.00001 s (seconds)
=====> Size (S) : 1000 B (Bytes)
=====> Model (M) : CM02
[0.000000] [simix_kernel/INFO] setting 'workstation_model' to 'compound'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'cpu_model' to 'Cas01'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'network_model' to 'CM02'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[S1:master:(1) 0.010010] [msg_test/INFO] Send completed (to C1). Transfer time: 0.010010
Aggregate bandwidth: 99900.099900
[S1:master:(1) 0.010010] [msg_test/INFO] Completed peer: C1 time: 0.010010
[C1:slave:(2) 0.010010] [msg_test/INFO] ==> Estimated Bw of FLOW[1] : 99900.099900 ;
message from S1 to C1 with remaining : 0.000000
=====<
=====> Bandwidth (B) : 1.000000e+05 B/s (Bytes per second)
=====> Latency (L) : 0.00001 s (seconds)
=====> Size (S) : 1000 B (Bytes)
=====> Model (M) : GTNets
[0.000000] [simix_kernel/INFO] setting 'workstation_model' to 'compound'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'cpu_model' to 'Cas01'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'network_model' to 'GTNets'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[S1:master:(1) 0.013140] [msg_test/INFO] Send completed (to C1). Transfer time: 0.013140
Aggregate bandwidth: 76103.500761
[S1:master:(1) 0.013140] [msg_test/INFO] Completed peer: C1 time: 0.013140
[C1:slave:(2) 0.013140] [msg_test/INFO] ==> Estimated Bw of FLOW[1] : 76103.500761 ;
message from S1 to C1 with remaining : 0.000000
=====<
=====> Bandwidth (B) : 1.000000e+05 B/s (Bytes per second)
=====> Latency (L) : 0.00001 s (seconds)
=====> Size (S) : 1000 B (Bytes)
=====> Model (M) : LegrandVelho
[0.000000] [simix_kernel/INFO] setting 'workstation_model' to 'compound'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'cpu_model' to 'Cas01'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'network_model' to 'LegrandVelho'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[S1:master:(1) 0.010974] [msg_test/INFO] Send completed (to C1). Transfer time: 0.010974
Aggregate bandwidth: 91128.086469
[S1:master:(1) 0.010974] [msg_test/INFO] Completed peer: C1 time: 0.010974
```

```
[C1:slave:(2) 0.010974] [msg_test/INFO] ==> Estimated Bw of FLOW[1] : 91128.086469 ;
message from S1 to C1 with remaining : 0.000000
=====><=====
```

3.2 Running the simulation in the Docker image

As the reader may have noted, the Docker image we produced only contains the binary code of the simulator but not the input files nor the perl script to run simulations. This is an intended separation of concerns and we believe it is a good practice to keep images as lightweight as possible and easier to maintain. We now describe how to rerun the simulation. We should first run the docker container.

```
1 docker run -ti aleggrand38/simgrid3_3_gtnets
```

Then the template XML input files should be copied within the container (\$CONTAINER corresponds to the container id of the container and is obtained either using `docker ps` or by querying the hostname within the container).

```
1 docker cp simutools09/instructions/01-onelink/onelink-d-template.xml \
2 $CONTAINER:/root/simutools09/01-onelink
3 docker cp simutools09/instructions/01-onelink/onelink-p-template.xml \
4 $CONTAINER:/root/simutools09/01-onelink
```

It is then possible to substitute the target parameters in these XML files and to run the simulation in the container:

```
1 cd /root/simutools09/01-onelink
2 sed -e s/bw/1.000000e+05/g -e s/lt/0.00001/g onelink-p-template.xml \
3 > /tmp/onelink-p.xml
4 sed -e s/size/1000/g onelink-d-template.xml > /tmp/onelink-d.xml
5 for model in CM02 GTNets LegrandVelho; do
6   echo ">=====<"
7   echo "=====> Model (M) : $model"
8   /root/simgrid-3.3/examples/msg/gtnets/gtnets \
9     /tmp/onelink-p.xml /tmp/onelink-d.xml \
10    --cfg=workstation_model:compound --cfg=cpu_model:Cas01 \
11    --cfg=network_model:$model;
12 done;
```

```
>=====<
=====> Model (M) : CM02
echo 'org_babel_sh_eoe'
echo 'org_babel_sh_eoe'
[0.000000] [simix_kernel/INFO] setting 'workstation_model' to 'compound'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'cpu_model' to 'Cas01'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'network_model' to 'CM02'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[S1:master:(1) 0.010010] [msg_test/INFO] Send completed (to C1). Transfer time: 0.010010
Aggregate bandwidth: 99900.099900
[S1:master:(1) 0.010010] [msg_test/INFO] Completed peer: C1 time: 0.010010
[C1:slave:(2) 0.010010] [msg_test/INFO] ==> Estimated Bw of FLOW[1] : 99900.099900 ;
message from S1 to C1 with remaining : 0.000000
=====<
=====> Model (M) : GTNets
[0.000000] [simix_kernel/INFO] setting 'workstation_model' to 'compound'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'cpu_model' to 'Cas01'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'network_model' to 'GTNets'
[0.000000] [xbt_cfg/INFO] type in variable = 2
<<<<=====>>>>
Dumping GTNETS topology information
== LINKID: 0
```

```

[Src] ID: 0, router?: 0, hosts[]: [ 0]
[Dst] ID: 1, router?: 0, hosts[]: [ 1]
>>>>=====<<<<
[S1:master:(1) 0.013140] [msg_test/INFO] Send completed (to C1). Transfer time: 0.013140
      Agregate bandwidth: 76103.500761
[S1:master:(1) 0.013140] [msg_test/INFO] Completed peer: C1 time: 0.013140
[C1:slave:(2) 0.013140] [msg_test/INFO] ==> Estimated Bw of FLOW[1] : 76103.500761 ;
      message from S1 to C1 with remaining : 0.000000
>=====
=====> Model      (M) : LegrandVelho
[0.000000] [simix_kernel/INFO] setting 'workstation_model' to 'compound'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'cpu_model' to 'Cas01'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[0.000000] [simix_kernel/INFO] setting 'network_model' to 'LegrandVelho'
[0.000000] [xbt_cfg/INFO] type in variable = 2
[S1:master:(1) 0.010974] [msg_test/INFO] Send completed (to C1). Transfer time: 0.010974
      Agregate bandwidth: 91128.086469
[S1:master:(1) 0.010974] [msg_test/INFO] Completed peer: C1 time: 0.010974
[C1:slave:(2) 0.010974] [msg_test/INFO] ==> Estimated Bw of FLOW[1] : 91128.086469 ;
      message from S1 to C1 with remaining : 0.000000

```

We could thus recover exactly the expected values which are reported on page 12.

3.3 Replicating the first series of simulations

Using the perl script should thus allow to re-execute the simulation. It comprises hard-coded absolute paths and a quick minor modification had thus to be made. Here is how to proceed:

- Within the container, we first create the directories that will host the simulation results:

```

1  mkdir -p /root/simutools09/01-onelink/bin
2  mkdir -p /root/simutools09/01-onelink/dat
3  mkdir -p /root/simutools09/01-onelink/log
4  mkdir -p /root/simutools09/01-onelink/tmp

```

- Then outside the container, we copy the template input files and simulation perl script:

```

1  docker cp simutools09/instructions/01-onelink/onelink-d-template.xml \
2  $CONTAINER:/root/simutools09/01-onelink
3  docker cp simutools09/instructions/01-onelink/onelink-p-template.xml \
4  $CONTAINER:/root/simutools09/01-onelink
5  docker cp simutools09/instructions/01-onelink/bin/sweep-parse.pl \
6  $CONTAINER:/root/simutools09/01-onelink/bin/sweep-parse.pl

```

- And finally back inside the container, we fix the absolute paths before running the simulations:

```

1  old_path="/home/velho/Development/projet-simgrid/simgrid/examples/msg/gtnets"
2  new_path="/root/simgrid-3.3/examples/msg/gtnets"
3  sed -i "s|$old_path|$new_path|g" /root/simutools09/01-onelink/bin/sweep-parse.pl
4  cd /root/simutools09/01-onelink/
5  ./bin/sweep-parse.pl sweep 1 1

```

```

Bandwidth array size is : 43
Changing working directory to /root/simgrid-3.3/examples/msg/gtnets
=====
Bandwidth (B) : 1.000000e+05 B/s (Bytes per second)
Latency (L) : 0.00001 s (seconds)
Size (S) : 1000 B (Bytes)
...

```

This worked like a charm! Unfortunately, according to the script, there are $40 \times 15 = 600$ (latency, bw) combinations, which run each for a bit more than a minute, hence about 10 hours solely for the first series of experiments. We did not let it run to the end but we checked that the parsing works and that all results matched for a hundred of combinations.

3.4 Running the analysis of the first series of experiments.

The analysis depends on master R script (`simutools09/instructions/01-oneLink/analyze.R`) which invokes perl and gnuplot. This is ugly but all pretty standard so we decided there was no need to rebuild a dedicated analysis environment and that it should run directly on our machine. To avoid messing up with the content of the original data, we decided to work in the `/tmp` of our machine as follows:

```
1 mkdir -p /tmp/simutools09/01-oneLink/dat/
2 mkdir -p /tmp/simutools09/01-oneLink/log/
3 mkdir -p /tmp/simutools09/01-oneLink/tmp/
4 mkdir -p /tmp/simutools09/01-oneLink/bin/
5 docker cp $CONTAINER:/root/simutools09/01-oneLink/dat/raw.data /tmp/simutools09/01-oneLink/dat/
6 cp simutools09/instructions/01-oneLink/analyze.R /tmp/simutools09/01-oneLink/
7 cp simutools09/instructions/01-oneLink/bin/* /tmp/simutools09/01-oneLink/bin/
```

```
1 source("analyze.R");
```

```
# Latency (SECONDS) Size (BYTES) Time (SECONDS)
Relax this may take some time
.....
.....
.....
.....
Candidates are X=0.934752791154703 and Y=10.6510810055123
The min is approximately: 0.0466609377572045
[1] "Hello!!!"
-Inf & NaN & NA \\
-Inf & NaN & NA \\
-Inf & NaN & NA \\
-Inf & NaN & NA \\
-Inf & NaN & NA \\
-Inf & NaN & NA \\
-Inf & NaN & NA \\
9.524 87.72 490.2 905.8 989.7 999 999.9-Inf & NaN & NA \\
-Inf & NaN & NA \\
-Inf & NaN & NA \\
-Inf & NaN & NA \\
"./tmp/gnuplotError.script" line 4: undefined variable: Inf
"./tmp/gnuplotError.script" line 4: undefined variable: Inf
There were 15 warnings (use warnings() to see them)
```

When running, a gnuplot window with a 3D graph popped up. There are error messages but the "Candidates are $X=0.934752791154703$ and $Y=10.6510810055123$ " message is really nice as these are the latency and bandwidth modifiers obtained through a custom linear regression and this is very familiar. The original paper reports .92 and 10.4 (page 5). The warnings and the differences come from the fact that the regression we just run was done using a smaller set of simulations because we didn't want to waste our time rerunning all the simulations.

4 Conclusion and take-away messages

Although we only replicated a fraction of the simulations conducted in the original article, they all perfectly match and we are confident that all the results would be reproduced with a few additional hours of efforts and enough time to run all the simulations

(several days actually). This is of little interest as GTNetS has been replaced in earlier versions of SimGrid by another packet level simulator: NS3.

We have shown in this article how to use modern tools such as the Docker Hub, the Debian snapshot archive, the Debuerreotype, GitHub, and Software Heritage. We have tried to demonstrate best practices and to highlight their effectiveness or potential shortcomings. Although they all require a relatively high level of operating system understanding and expertise, we believe they are all now mature enough and sufficiently easy to use both for such kind of computer "archaeology" and for a daily usage in a research context (, which greatly eases the task of anyone trying to reuse or reproduce the work). A sound question to ask is: "Would anyone other than the original authors have succeeded in reproducing this work?". A fair answer is probably no.

- First, three different archives were needed: the first one was easy to find, the second one was publicly available but deeply hidden so it is unlikely anyone else than the original authors would have found it, and although the third one was also available on the Internet, it was not very visible and we initially recovered from the hard drive of one of the two original authors.
- Second, even after gathering the three archives, rebuilding the software environment, correctly linking and running the simulation was possible but required such a good amount of faith that we believe anyone else than the original authors would have easily given up.

At the time of writing of the original article, Pedro Velho had put a significant effort in documenting the whole workflow and relying on standard tools such as R, perl, and make to automate as much work as possible. Yet, we made the three following mistakes from a reproducibility perspective:

1. We never reached the point where a full automation was done and delegated to a third party. In particular as we had no satisfying nor standard tool to distribute the workload on a cluster. So we kept track of simulation outputs and intermediate results manually. This good organization has been a life saver when trying to reproduce and check the results. If we had to redo such work today, we would probably use something like `snakemake` and `org-mode` notebooks to easily automate and document the whole work.
2. We underestimated the URL rot effect. Although all our work was version controlled, moving from a development platform to another made information and archive retrieval more difficult than we anticipate. Although all the archives were finally available, it took us an inordinate amount of time to locate them.. Cleaning up is rarely done after publishing, hence the need to do it on the fly. It turns out that Pedro Velho had taken care to clean and to make all the data he had produced during his PhD thesis available on his webpage. The policy in our lab is to maintain the webpage of former members so all the data is still available but not easily found. Using a perennial archive such as Zenodo would be the recommended way to proceed nowadays but this archive did not exist by then.
3. Finally, we underestimated the importance of capturing every information on software environment. A few ones related to processor architecture and compilers were available but it was lacunar. Fortunately, we only relied on standard open source software and from the dates, it was not too hard to identify which software must have been used and we have been able to rebuild a functional software environment at low cost, solely from binary packages. Controlling this environment and making it easily available and usable is definitely the way to go with tools like Docker but this lightweight virtualization was not as easy to use back then.

Overall this reproducibility challenge was an excellent experience to face the effect of time even on a relatively short (10 years) time period. It was also very positive to realize that over the last decade, several very good tools and practices have emerged to address exactly the difficulties we faced back then (workflows and notebooks to handle computations, software and data archives to fight against link rot, container and stable packaging systems to manage software environments).

References

1. P. Velho and A. Legrand. "Accuracy Study and Improvement of Network Simulation in the SimGrid Framework." In: **SIMUTools'09, 2nd International Conference on Simulation Tools and Techniques**. Rome, Italy, Mar. 2009.
2. K. Fujiwara and H. Casanova. "Speed and Accuracy of Network Simulation in the SimGrid Framework." In: **Intl. Conf. on Performance Evaluation Methodologies and Tools**. Nantes, France, 2007, 12:1–12:10.
3. S. Naicken, A. Basu, B. Livingston, and S. Rodhetbhai. "Towards Yet Another Peer-to-Peer Simulator." In: **Proc. Fourth International Working Conference Performance Modelling and Evaluation of Heterogeneous Networks (HET-NETs '06)**. Sept. 2006.
4. H. Casanova and L. Marchal. **A Network Model for Simulation of Grid Application**. Tech. rep. 2002-40. LIP, Oct. 2002.
5. L. Massoulié and J. Roberts. "Bandwidth Sharing: Objectives and Algorithms." In: **INFOCOM**. 1999, pp. 1395–1403.
6. G. Marfia, C. Palazzi, G. Pau, M. Gerla, M. Sanadidi, and M. Roccetti. "TCP Libra: Exploring RTT-Fairness for TCP." In: **NETWORKING 2007. Ad Hoc and Sensor Networks, Wireless Networks, Next Generation Internet**. Ed. by I. F. Akyildiz, R. Sivakumar, E. Ekici, J. C. d. Oliveira, and J. McNair. Vol. 4479. Lecture Notes in Computer Science. Springer Berlin Heidelberg, 2007, pp. 1005–1013.
7. P. Velho, L. Schnorr, H. Casanova, and A. Legrand. "On the Validity of Flow-level TCP Network Models for Grid and Cloud Simulations." In: **ACM Transactions on Modeling and Computer Simulation** 23.4 (Oct. 2013).
8. P. A. Madeira de Campos Velho. "Accurate and Fast Simulations of Large-Scale Distributed Computing Systems." Theses. Université Grenoble Alpes, July 2011.
9. G. F. Riley. "The Georgia Tech Network Simulator." In: **ACM SIGCOMM workshop on Models, Methods and Tools for Reproducible Network Research**. Karlsruhe, Germany, 2003, pp. 5–12.
10. A. Medina, A. Lakhina, I. Matta, and J. Byers. **BRITE: Universal Topology Generation from a User's Perspective**. Available at <https://www.cs.bu.edu/brite/publications/usermanual.pdf>. Apr. 2001.
11. B. M. Waxman. "Routing of Multipoint Connections." In: **IEEE Journal on Selected Areas in Communications** 6.9 (Dec. 1988), pp. 1617–1622.